

CMPUT 291 Project 2 Report

By Omar Almokdad, Gemma Marcinkoski, Wai Yi Low

Implementation Techniques & Analysis

B-TREE & HASH:

The implementation of the B-tree file, and the hash file was basically identical in terms of the navigation, and algorithms used to retrieve the data, the only change was in the way the database was opened to write to or parse. The B-tree file called the opener function in the bsddb3 library with the B_TREE flag, while the Hash file called the function with the HASH tag. This tag is what tells the database how to structure the data.

The BTREE and HASH databases gave similar response times to each other in the three different scenarios: searching by key, searching by value, and searching by range with lower and upper limit, as seen in the tables below.

Our first method, searching by key, generated the results from both databases is in average of about 200 microseconds with insignificant variance. Searching by key took the shortest time among the three scenarios because it uses the key, which is the main factor in how the data is structured.

For the other two methods, searching by value and searching by range, generated results in average of over 100,000 microseconds. That is because a sequential search is required for both of them, while the data is not structured sequentially

INDEXFILE:

The implementation of the index file was a little bit different than the other two. For the index file, as the database was created, a second database was created along-side it, with the values set up to have their own index with a B-Tree structure. This allows both searching by key, and by value to be more efficient, and avoids sequentially searching the database in these two cases.

For the method for searching by key, the implementation is exactly the same as the B-Tree implementation, since it is very efficient on its own, averaging at about 200 microseconds.

However, for the search by value, the index file steps in and is treated as a stand-alone database to retrieve the key. This is efficient because the index file structure is based on the values being the new keys, thus giving the full efficiency of a B-Tree, averaging at about 200 microseconds as well.

A big drop is seen in the time needed to do a ranged search as well, because shortcuts to the searching were implemented to speed it up. In the Indexfile implementation, according to the cursor, the keys are indexed alphabetically, and in that case, the search only has to start from the lower bound until it reached the upper bound, instead of searching the whole database, this drops the range search average from the 100,000's to the 10,000's of milliseconds.

Experimental Results (Time taken to execute in milliseconds)

| B-Tree | Key Search | Data Search | Range Search | # of Entries in Range |
|---------------|------------|-------------|--------------|-----------------------|
| Test 1 | 159 | 109851 | 114220 | 3776 |
| Test 2 | 338 | 103225 | 122203 | 3830 |
| Test 3 | 183 | 103806 | 109262 | 3900 |
| Test 4 | 269 | 112935 | 110991 | 3773 |
| Average | 237.25 | 107454.3 | 114169 | 3819.75 |

| Hash | Key Search | Data Search | Range Search | # of Entries in Range |
|-------------|------------|-------------|--------------|-----------------------|
| Test 1 | 183 | 119030 | 136416 | 3776 |
| Test 2 | 177 | 121662 | 133407 | 3830 |
| Test 3 | 180 | 122166 | 131969 | 3900 |
| Test 4 | 314 | 128582 | 130380 | 3773 |
| Average | 213.5 | 122860 | 133043 | 3819.75 |

| Index | Key Search | Data Search | Range Search | # of Entries in Range |
|--------------|------------|-------------|--------------|-----------------------|
| Test 1 | 204 | 196 | 12334 | 3776 |
| Test 2 | 308 | 195 | 11934 | 3830 |
| Test 3 | 193 | 323 | 13095 | 3900 |
| Test 4 | 314 | 200 | 11050 | 3773 |
| Average | 254.75 | 228.5 | 12103.25 | 3819.75 |

Summary of Results

| Queries | B-Tree | Hash Table | Index File |
|--------------|--------|------------|------------|
| Key Search | 237 | 214 | 255 |
| Data Search | 107454 | 122860 | 229 |
| Range Search | 114169 | 133043 | 12103 |